



JOHNS HOPKINS
CAREY BUSINESS SCHOOL

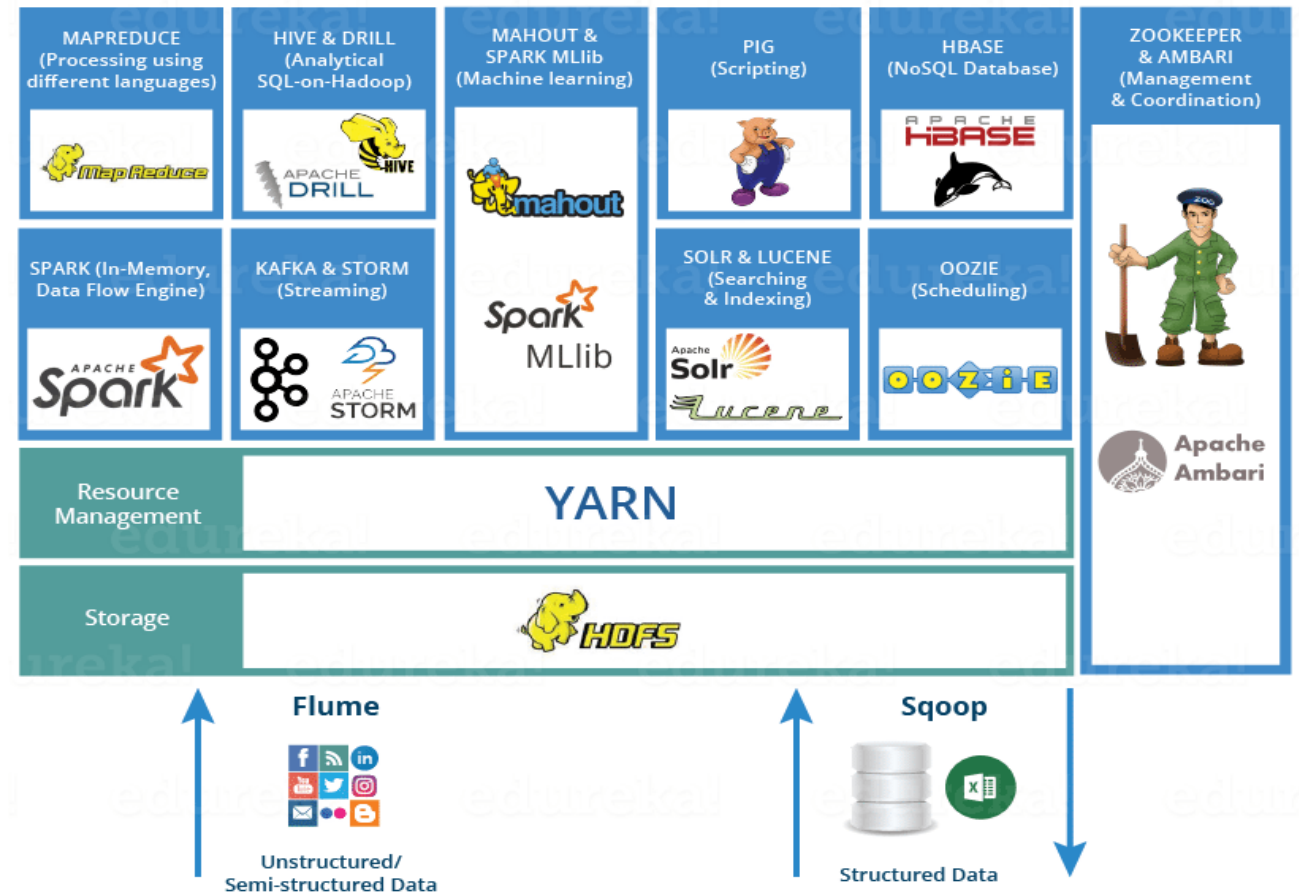
HDFS and YARN

BU.330.740 Large Scale Computing on the Cloud

Minghong Xu, PhD.
Associate Professor

Hadoop Overview

- » An open source framework for writing and running distributed applications that process large amount of data
- » Google: first to publicize MapReduce for scaled data processing
- » Doug Cutting: develop the first version of Hadoop





Key Components

- » Distributed File System: HDFS
 - Most low-level knowledge of this course
- » Operating System: YARN
- » Original Distributed Processing Engine: MapReduce
- » Improved MapReduce: Spark
- » Distributed Query Language: Hive
- » Distributed Scripting Language: Apache Pig



Distributed File System



Hadoop Distributed File System

» **Problem 1:** *Data is too big to store on one machine.*

» **Solution:** Store the data on multiple machines!

» **Problem 2:** *Very high end machines are too expensive*

» **Solution:** Run on commodity hardware!

» **Problem 3:** *Commodity hardware can fail*

» **Solution:** Software is intelligent enough to handle hardware failure!

» **Problem 4:** *What happens to the data if the machine storing the data fails?*

» **Solution:** Replicate the data!



Solutions (Cont.)

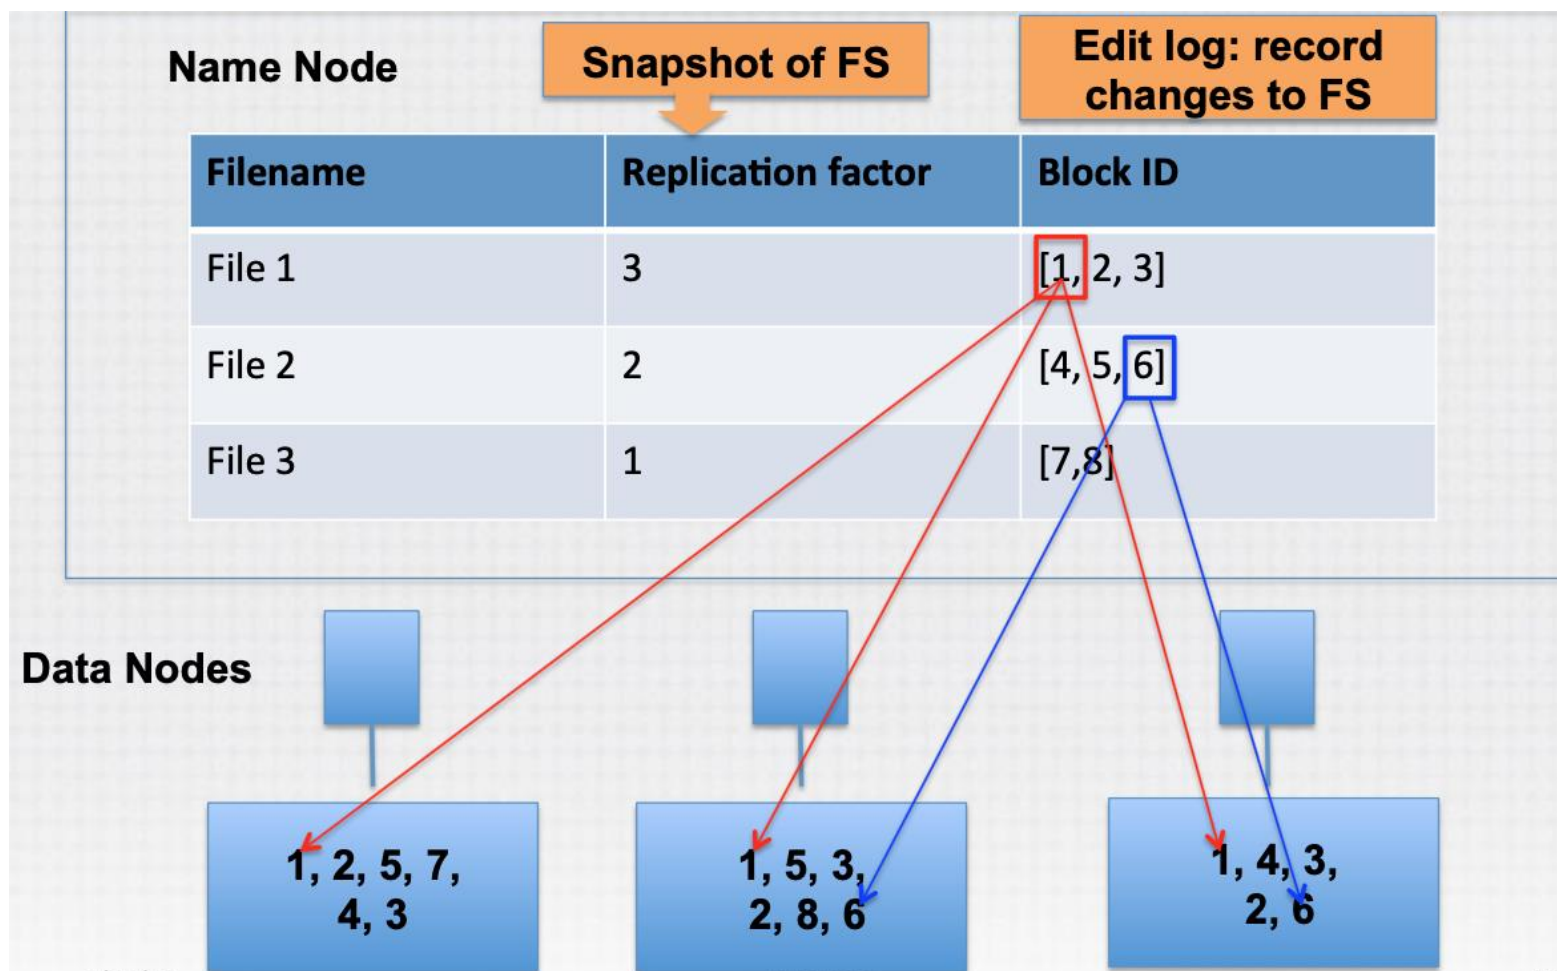
- » **Problem 5:** *How can distributed machines organize the data in a coordinated way?*
- » Insights from your own team work experience...
 - *How to assign tasks to different workers in an efficient way?*
 - *What happens if tasks fail?*
 - *How do workers exchange results?*
 - *How to synchronize distributed tasks allocated to different workers?*
- » **Solution:** primary-secondary structure



HDFS Primary-Secondary Architecture

- » Single NameNode
 - Sometimes a backup: secondary NameNode
- » Many (Thousands) DataNodes
- » Files are split into fixed sized blocks and stored on data nodes
- » Data blocks are replicated for fault tolerance and fast access
 - By default: 3

Illustration



Adapt from K. Zhang's notes, Spring 2019



NameNode

- » Manages file system namespace, and file metadata
- » Mapping file to list of blocks
- » Mapping of datanode to list of blocks
- » Monitor datanode health
- » Replicate missing blocks


DataNode



- » Handle block storage & block integrity
- » Periodically send reports to NameNode
- » Clients access the blocks directly from data nodes
- » *Q: Why not access blocks through NN?*
- » Reasons:
 - Prevent NN from being the bottleneck of the cluster
 - Allow HDFS to scale to large number of concurrent clients
 - Spread the data traffic across the cluster

Data Replicate

- » **Frist replica** is put on one node in the local rack
- » **Second one** is put on a node in a different (remote) rack
- » **Third one** is on a different node in the same remote rack
- » **Additional** replicas are randomly placed
- » Objectives: load balancing, fast access, fault tolerance.





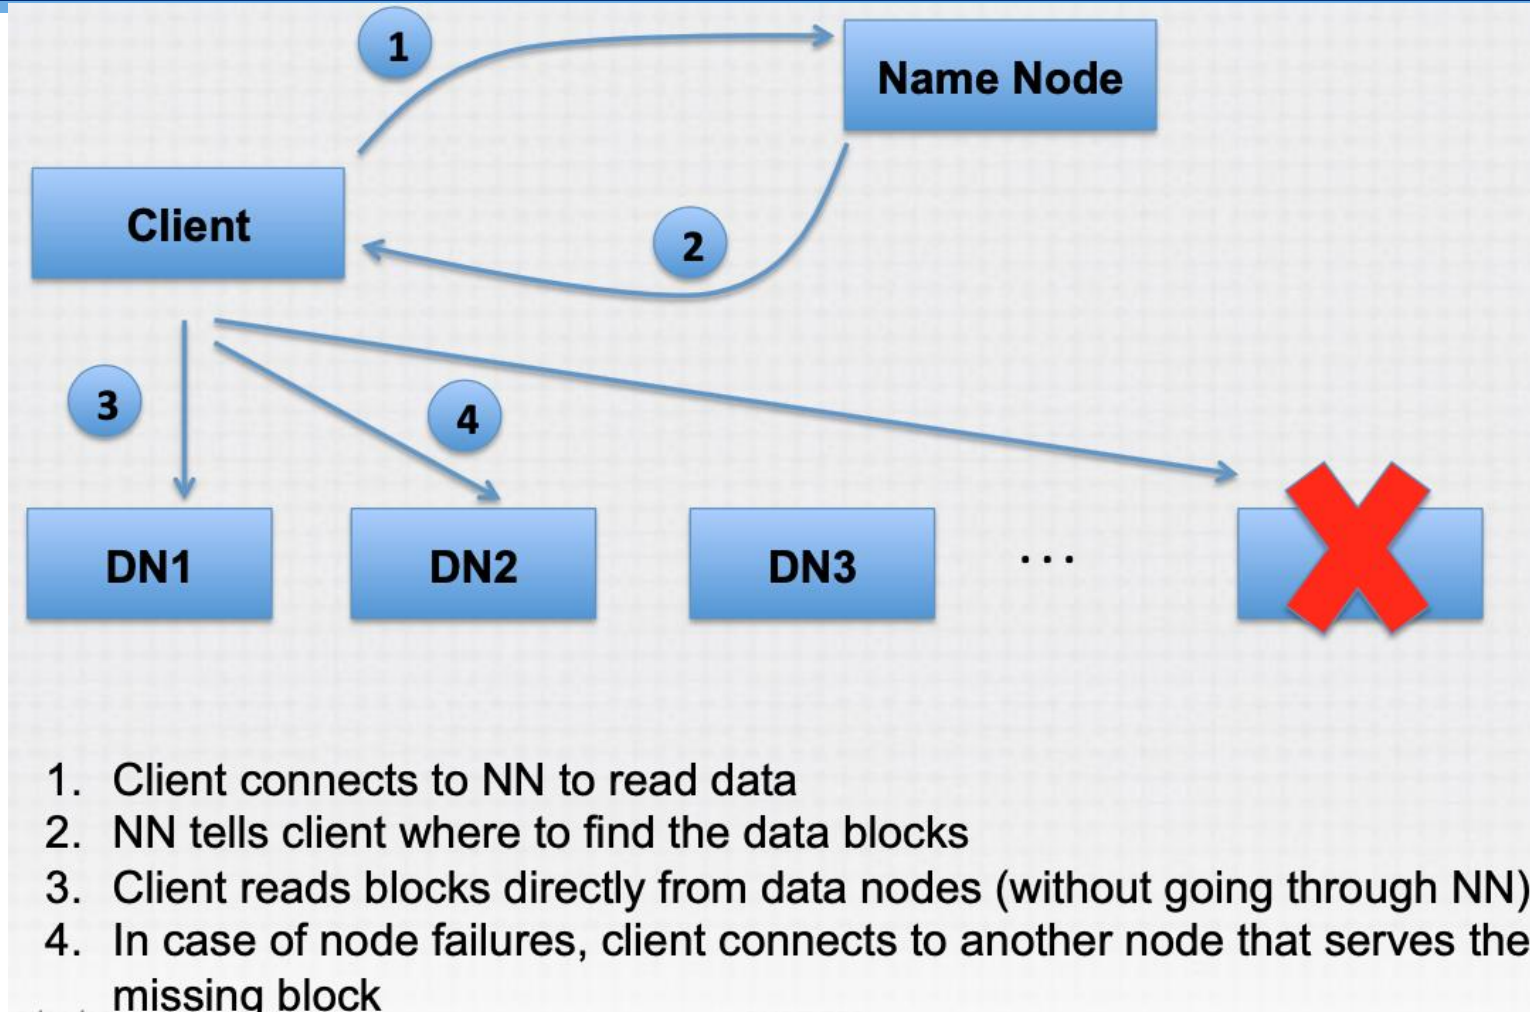
HDFS Network Topology

- » The critical resource in HDFS is **bandwidth**, distance is defined based on that
- » Measuring bandwidths between any pair of nodes is too complex
- » **Basic Idea:**
 - Processes on the same node
 - Different nodes on the same rack
 - Nodes on different racks in the same data center (cluster)
 - Nodes in different data centers



**Bandwidth
becomes less**

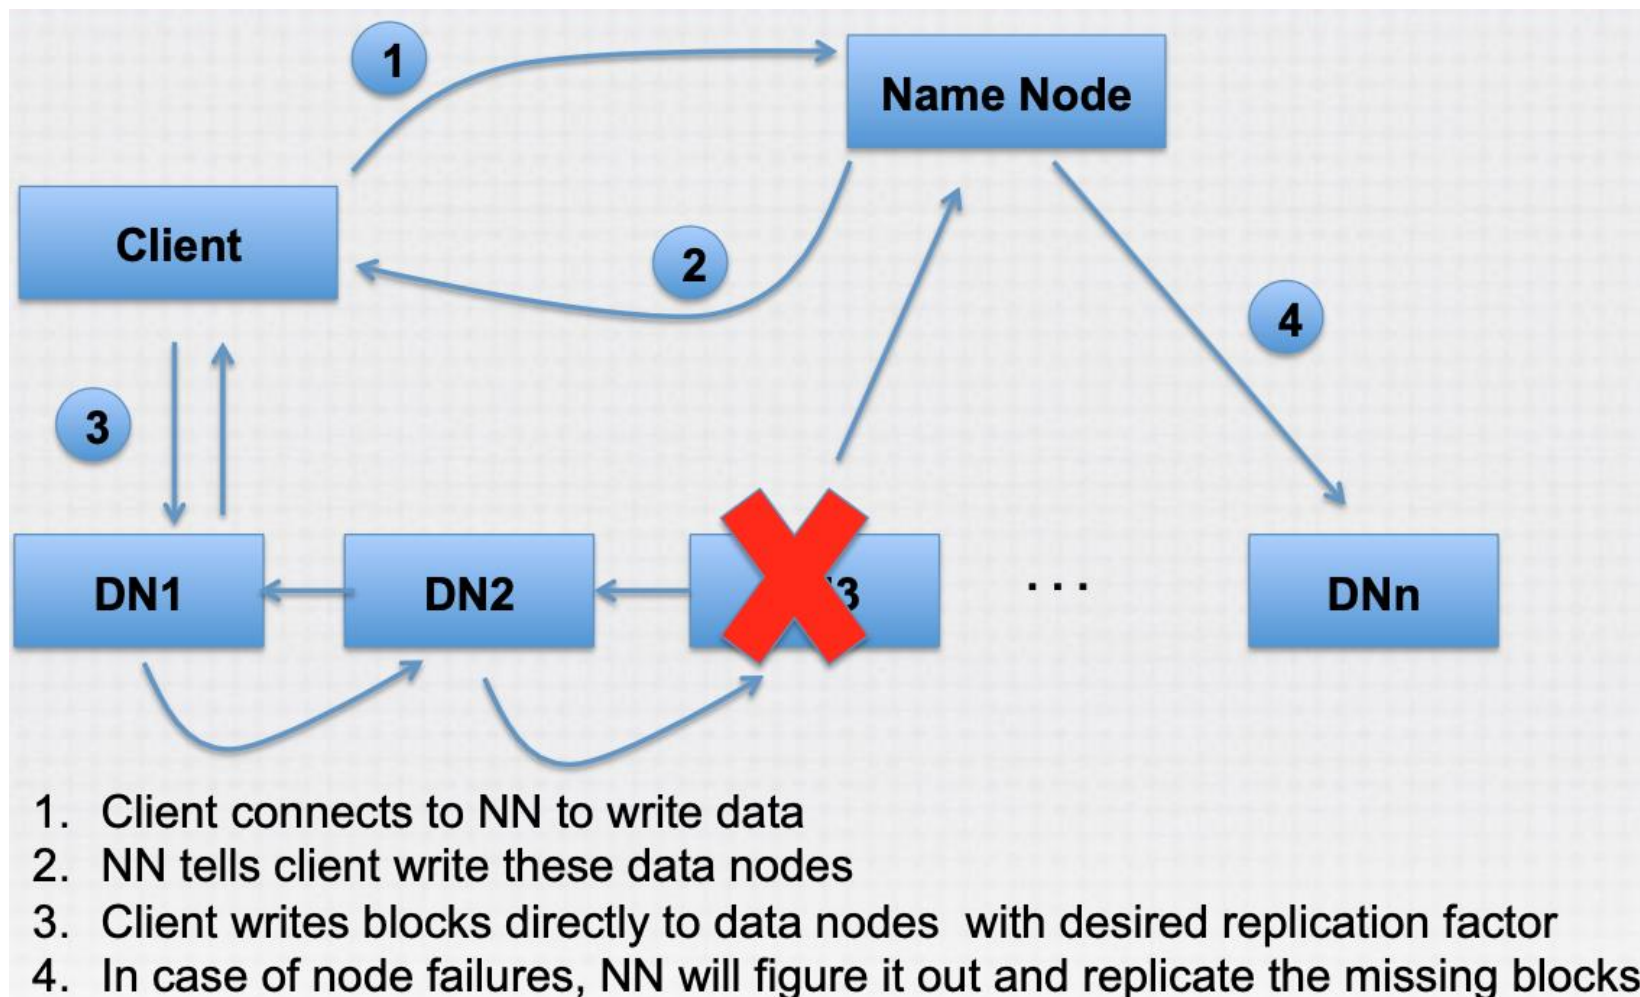
Data Read



1. Client connects to NN to read data
2. NN tells client where to find the data blocks
3. Client reads blocks directly from data nodes (without going through NN)
4. In case of node failures, client connects to another node that serves the missing block

Adapt from K. Zhang's notes, Spring 2019

Data Write

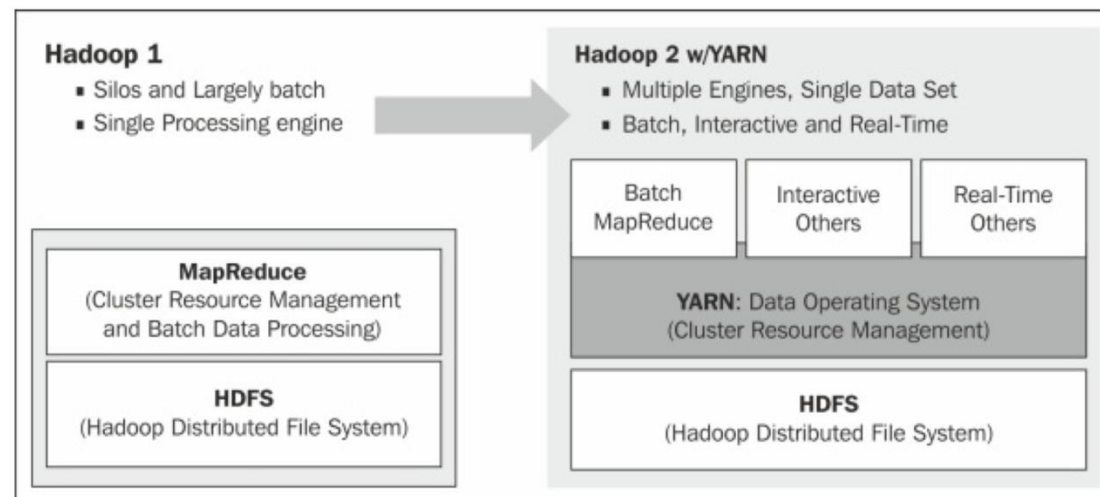




Yet Another Resource Negotiator

What is YARN?

- » Cluster resource management system for Hadoop
- » Introduced in Hadoop 2 to improve MapReduce implementation
- » Connect between high level applications (Spark, HBase) and low level Hadoop environment
- » Large-scale, distributed operating system for big data applications



Components



» Resource Manager

- 1 per cluster
- Track resource in a cluster, schedule applications
- Single point of failure, but can be restored in case of failures

» Application Master

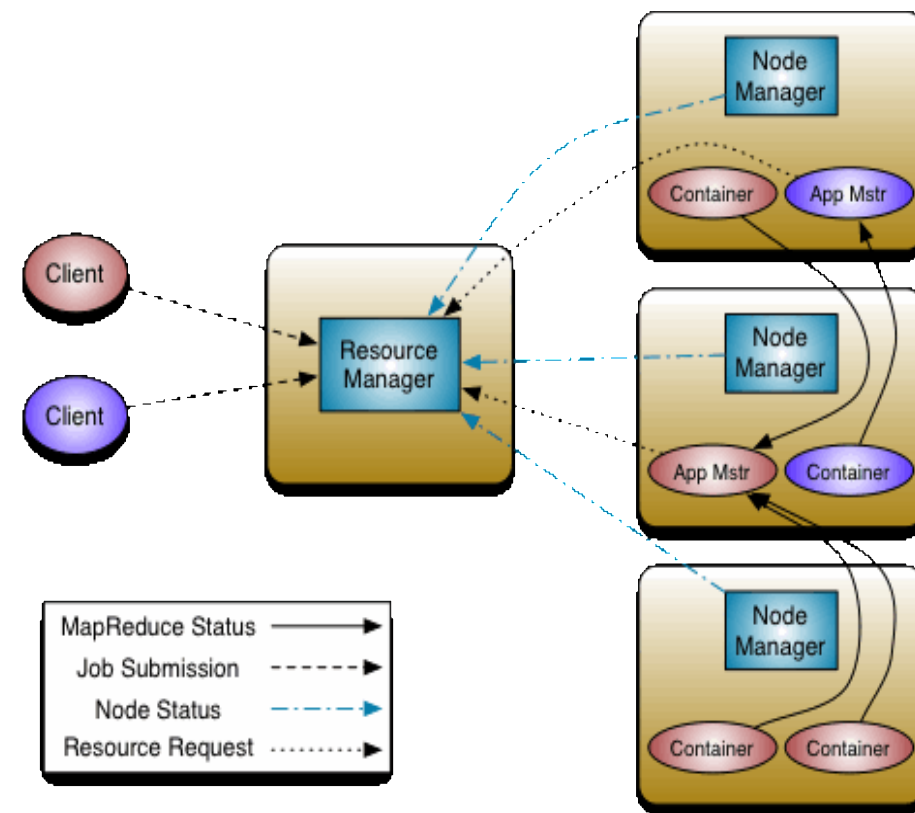
- Run in a separate process on a datanode
- One instance per application
- Send status and resource needs to RM
- Can run lightweight task on the same node

» Node Manager

- 1 per node
- Monitor node resources such as CPU, Memory, Disk space, Network etc.
- Collect log data, report to Resource Manager

Steps to run Yarn application

- » Client make request to Resource Manager to run application
- » Resource Manager request Node Manager to allocate container for creating Application Master instance on available node
 - Container: basic unit of hardware allocation
- » When Application Master instance already run, it sends request (heartbeat, resource needs) to Resource Manager





Scheduling in Yarn

- » Important task since Hadoop cluster is shared between many users and tasks; which to run first?
- » Hadoop 1: FIFO scheduler with fixed cpu, memory, disk count
- » Hadoop 2: allow Capacity and Fair schedulers with dynamic cpu, memory and disk count
- » Yarn supports all three

Illustrations

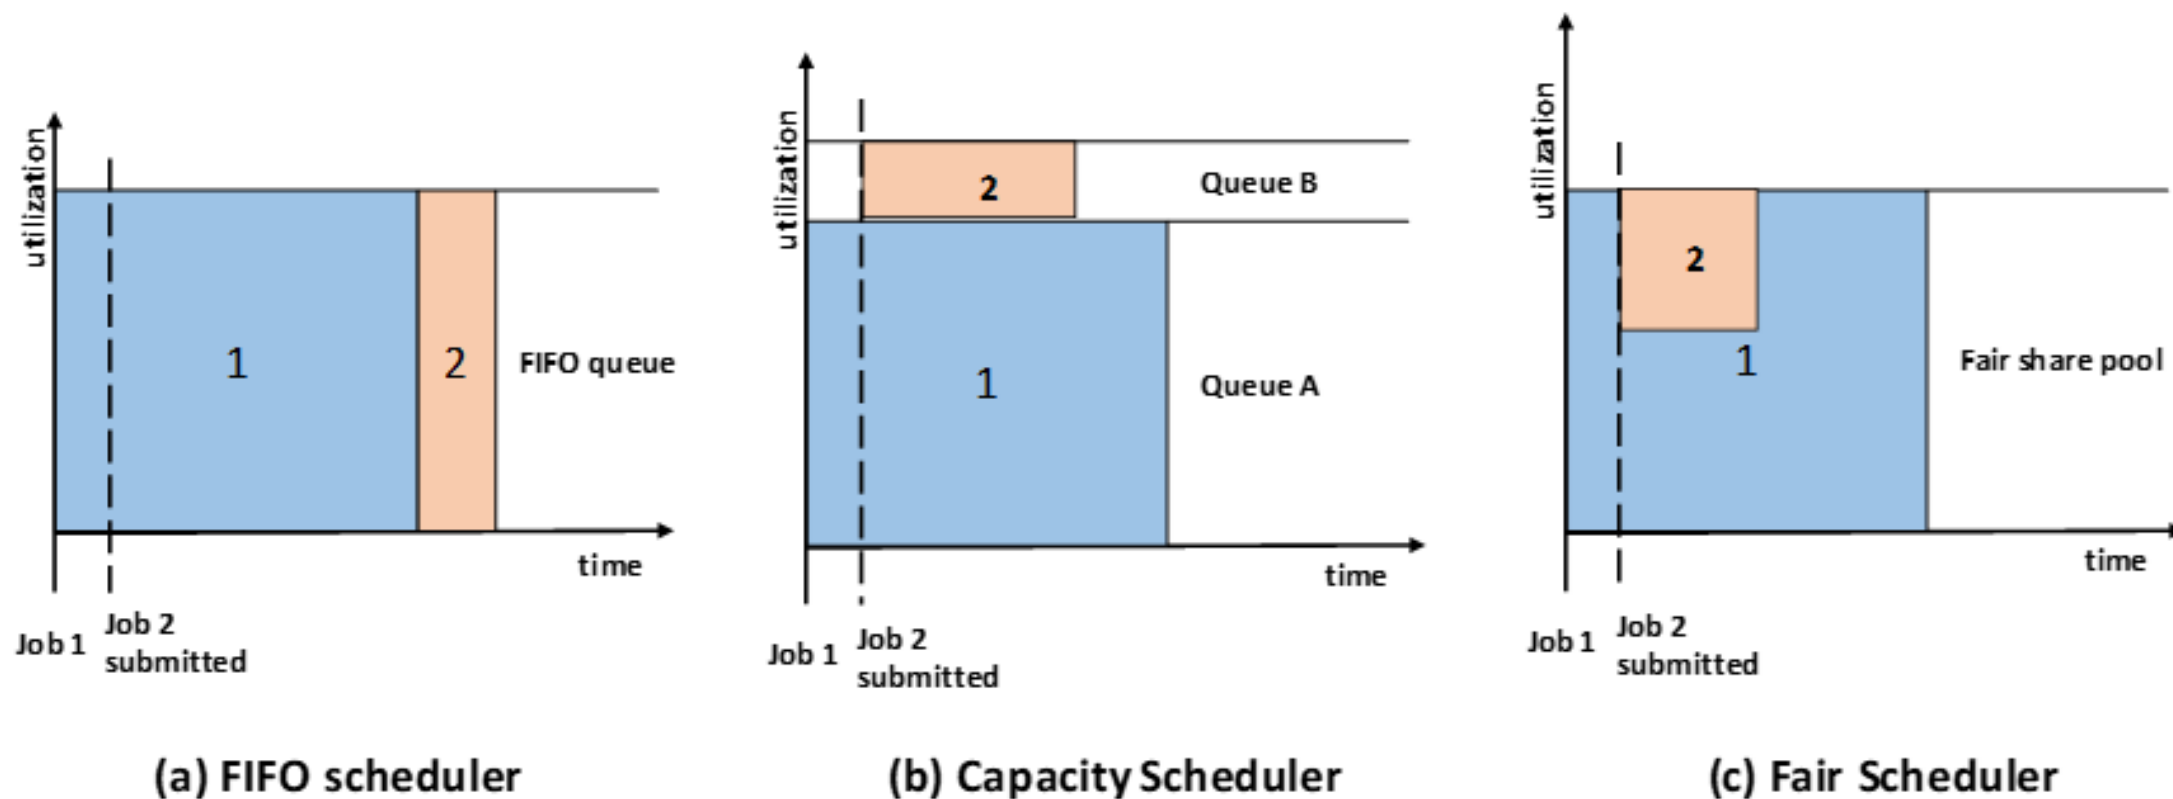


Figure 1: YARN Schedulers' cluster utilization vs. time